

VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



Introduction to A.I (CO3061)

Assignment 2

Game playing with Search Algorithms

Advisor(s):	Tran Hong Tai	
Student(s):	Huynh Van Anh Hoangi	2252229
	Vu Quoc Viet	2353326
	Tran Xuan Hao	2252191
	Tran Gia Huy	2252264
	Vu Song Anh	2252048

Hồ Chí Minh city, April 12th 2026



Table of contents

1	Overview	3
1.1	Introduction	3
1.2	Objectives	3
1.3	Scope	3
2	Adversarial Search: Alpha-Beta Pruning	4
2.1	State-Space Model for Chess	4
2.2	Negamax Symmetry and Score Inversion	4
2.3	Alpha-Beta Pruning	4
2.4	Quiescence Search	5
2.4.1	The Horizon Effect Challenge	5
2.4.2	Implementation Strategy	5
2.4.3	Move Ordering in Quiescence	6
2.5	Move Ordering (MVV-LVA)	6
2.5.1	Heuristic Logic and Formula	6
2.5.2	Code Implementation	7
2.5.3	Impact on Performance	7
2.6	Performance Optimization: Zobrist Hashing	7
2.6.1	Code Implementation: Zobrist and Cache	8
3	Evaluation Function	8
3.1	Heuristic Overview	8
3.2	Material Weighting	9
3.3	Piece-Square Tables (PST)	9
3.3.1	Positional Strategy	9
3.3.2	Code Implementation	10
3.3.3	Evaluation Formula	10
3.4	Mobility and Penalties	10
3.4.1	Advanced Heuristic Factors	10
3.4.2	Combined Evaluation Sum	11
4	Experimental Evaluation	11
4.1	Agent Skill Levels	11
4.2	Experimental Setup	12
4.3	Results vs. Random Agent	12
4.4	Self-Play Analysis	13
5	References	13



1 Overview

1.1 Introduction

Adversarial search algorithms are fundamental techniques in Artificial Intelligence (AI) that are used to explore a competitive game space in order to identify optimal moves that maximize an agent's utility while accounting for the opponent's strategy.

In this project, we investigate the application of these techniques to the game of Chess, a classic adversarial game with a vast state space that presents a significant challenge for classical search algorithms.

Both players in Chess aim to reach a terminal goal state—checkmate—while navigating a deep search tree where every action corresponds to a valid move according to international rules. By exploring the state space using informed strategies, a search algorithm attempts to find configurations that satisfy all constraints and therefore represent a valid path to victory.

The primary goal of this project is to implement an agent using the Negamax algorithm with Alpha-Beta pruning to analyze how heuristic functions can significantly improve search efficiency in high-complexity environments.

1.2 Objectives

The main objectives of this assignment are:

1. To understand the concept of adversarial state-space search in AI and how complex games like Chess can be modeled as search problems.
2. To formulate Chess as a state-space search problem, including the definition of piece layouts, legal actions, and win/draw conditions.
3. To implement an intelligent game-playing agent using classical searching algorithms, specifically excluding Machine Learning.
4. To design a robust heuristic evaluation function to guide the agent through deeper levels of the search tree.
5. To establish ten distinct skill levels for the agent and evaluate their performance against a random baseline.

1.3 Scope

The scope of this project focuses on applying classical adversarial search algorithms to solve the game of Chess.

The study is conducted within a controlled experimental environment and includes the following tasks:

- Modeling the game of Chess as a state-space search problem, including the formal definition of board states, castling rights, and transition rules.
- Implementing the Negamax search strategy optimized with Alpha-Beta pruning and Quiescence search to handle the large branching factor of Chess.
- Developing the complete solution in Python 3 using modern dependency management tools like uv.

- Creating multiple skill profiles ranging from Level 1 to Level 10 to evaluate algorithm performance under varying depth and node constraints.
- Conducting experiments to measure agent efficiency using metrics such as node expansion, win rate against random opponents, and execution time.
- Presenting the methodology, implementation details, experimental results, and analysis in a comprehensive technical report.

The project concentrates on demonstrating the fundamental concepts of adversarial search algorithms and their practical application to chess-playing solving.

2 Adversarial Search: Alpha-Beta Pruning

2.1 State-Space Model for Chess

The chess problem is modeled as a state-space search where each state represents a specific configuration of the 8×8 board.

- **Initial State:** The standard starting position of 32 pieces arranged on the board .
- **Actions:** A set of all legal moves including specialized rules such as castling, en passant, and pawn promotion.
- **Transition Model:** Applying a move results in a new board state and toggles the turn to the opponent.
- **Terminal Test:** States where the game ends via checkmate, stalemate, or insufficient material.

2.2 Negamax Symmetry and Score Inversion

The core of the search implementation lies in the recursive symmetry of the Negamax algorithm. Instead of using two different logic branches for White and Black, we use a single mathematical framework:

$$\text{Score}_{\text{player}} = -\text{Negamax}(\text{depth} - 1, -\beta, -\alpha) \quad (1)$$

In our implementation within `agent.py`, this is captured by the following code snippet:

```
1 score = -self._negamax(depth - 1, -beta, -alpha)
```

Source Code 1: Negamax recursive call in `agent.py`

This inversion works because in a zero-sum game like Chess, a gain of +100 for White is exactly a loss of -100 for Black. By negating the value returned from the deeper search level, the current player always treats the value as "their own" score to maximize, regardless of whether they are playing as White or Black.

2.3 Alpha-Beta Pruning

To manage the massive branching factor of chess, **Alpha-Beta Pruning** is integrated into the search process.

- **Alpha (α):** Represents the minimum score the maximizing player is assured of.

- **Beta (β):** Represents the maximum score the minimizing player is assured of.

The search prunes branches as soon as it finds a move that proves worse than a previously examined option, significantly reducing the number of nodes expanded without affecting the final decision.




Figure 1: Alpha-Beta Pruning: Nodes N4 and N5 are maintained in the diagram to show the potential paths, but are marked with red double-slashes to indicate they were pruned by the `break` condition in `agent.py`.

2.4 Quiescence Search

2.4.1 The Horizon Effect Challenge

A major limitation of fixed-depth search algorithms is the *horizon effect*. This occurs when the search ends at a depth where the board state is volatile (e.g., in the middle of a piece exchange). The static evaluation function `material_and_pst` may return a highly optimistic score because it cannot see that a counter-capture or checkmate occurs on the very next move beyond the "horizon".

2.4.2 Implementation Strategy

To stabilize the evaluation, we implemented the `_quiescence` function in `agent.py`. Instead of returning the static score at `depth == 0`, the search is extended specifically for "noisy" moves:

- **Standing Pat:** The agent first calculates a "standing pat" score using the static evaluation. If this score is already high enough to cause a beta-cutoff, the search returns immediately to save time.
- **Selective Move Set:** Only capture moves are considered during this extension. This prevents an exponential explosion of the search tree while ensuring tactical blunders are caught.
- **Search Depth Management:** As defined in the `SkillProfile`, higher levels (7-10) are assigned a larger `quiescence_max_depth` (up to 4 plys) to handle complex tactical sequences.



Figure 2: Comparison: Standard search stops at Depth 4 and misses a capture, while Quiescence Search extends to Depth 5 to reach a 'quiet' and accurate state.

2.4.3 Move Ordering in Quiescence

Even within the quiescence search, we apply **MVV-LVA** ordering to ensure that the most promising captures (e.g., a Pawn taking a Queen) are evaluated first, further optimizing the pruning efficiency of the extra search plies.

2.5 Move Ordering (MVV-LVA)

The efficiency of Alpha-Beta pruning is highly dependent on the order in which moves are explored. If the best moves are evaluated early, more branches can be pruned, significantly reducing the search space [?, ?]. To achieve this, our agent implements the **MVV-LVA** (Most Valuable Victim - Least Valuable Attacker) heuristic within the `order_moves` function.

2.5.1 Heuristic Logic and Formula

The core idea is to prioritize captures where a high-value piece (Victim) is taken by a low-value piece (Attacker). Each move is assigned a priority score based on the piece values defined in our system:

```
1 def _mvv_lva(game: ChessGame, m: Move) -> int:
2     victim = 0
3     if game.board[m.to_r][m.to_c]:
4         victim = PIECE_VALUE[game.board[m.to_r][m.to_c][1]]
```

```
5 elif game.ep_square and (m.to_r, m.to_c) == game.ep_square:
6     victim = PIECE_VALUE[PAWN]
7
8     attacker = PIECE_VALUE[game.board[m.from_r][m.from_c][1]]
9     return victim * 100 - attacker
```

Source Code 2: MVV-LVA Implementation in agent.py

- **Material Values:** Pawn: 100, Knight: 320, Bishop: 330, Rook: 500, Queen: 900.
- **Scoring Formula:** As implemented in the `_mvv_lva` function, the priority is calculated as:

$$Score = Value_{Victim} \times 100 - Value_{Attacker} \quad (2)$$

The multiplier of 100 ensures that the value of the victim is the primary sorting factor, while the attacker's value acts as a tie-breaker to favor lower-value attackers.

2.5.2 Code Implementation

The following snippet from `agent.py` illustrates how this heuristic is used to sort the move list before the Negamax search begins:

```
1 def _mvv_lva(game: ChessGame, m: Move) -> int:
2     victim = 0
3     if game.board[m.to_r][m.to_c]:
4         victim = PIECE_VALUE[game.board[m.to_r][m.to_c][1]]
5     elif game.ep_square and (m.to_r, m.to_c) == game.ep_square:
6         victim = PIECE_VALUE[PAWN]
7
8     attacker = PIECE_VALUE[game.board[m.from_r][m.from_c][1]]
9     return victim * 100 - attacker
```

Source Code 3: MVV-LVA Implementation in agent.py

2.5.3 Impact on Performance

By examining promising captures first, the Alpha-Beta algorithm triggers cut-offs much earlier in the search tree. This optimization is critical for maintaining a competitive search depth (up to 7-8 plys) within the assigned 0.35s time limit per move.

2.6 Performance Optimization: Zobrist Hashing

To further enhance search efficiency, the agent implements Zobrist Hashing to uniquely identify board configurations. This technique allows the system to recognize transpositions—positions reached through different sequences of moves.

- **Evaluation Cache:** We implement a Transposition Table using an LRU (Least Recently Used) cache policy via the `_EvalLRU` class. Static evaluation results are stored and keyed by a 64-bit Zobrist hash, with a capacity of 80,000 entries to balance memory and speed.
- **Transposition Handling:** When the search encounters a position it has already evaluated, it retrieves the score instantly. This is especially effective in endgames where the same position is reached frequently, allowing for a much deeper search within the same time budget.
- **Zobrist Keys and XOR Logic:** Each piece-square combination is assigned a unique random 64-bit integer. The board hash is updated efficiently using bitwise XOR operations.

2.6.1 Code Implementation: Zobrist and Cache

The following snippet from `agent.py` illustrates the initialization of Zobrist keys and the implementation of the LRU evaluation cache:

```
1 _ZR = random.Random(0x5E7A8C)
2 _ZOB_SQ: List[List[List[int]]] = [
3     [[_ZR.getrandbits(64) for _ in range(12)] for _ in range(8)] for _ in range(
4         8)
5 ]
6
7 def _zobrist_pieces(game: ChessGame) -> int:
8     h = 0
9     for r in range(8):
10         row = game.board[r]
11         for c in range(8):
12             p = row[c]
13             if p:
14                 h ^= _ZOB_SQ[r][c][_piece_z_idx(p)]
15     return h
16
17
18 class _EvalLRU:
19     """Static eval cache keyed by piece layout (Zobrist); big win in deep search
20     """
21     __slots__ = ("_d", "_cap")
22
23     def __init__(self, cap: int = 80_000) -> None:
24         self._d: "OrderedDict[int, int]" = OrderedDict()
25         self._cap = cap
26
27     def get(self, key: int) -> Optional[int]:
28         v = self._d.get(key)
29         if v is None:
30             return None
31         self._d.move_to_end(key)
32         return v
33
34     def put(self, key: int, val: int) -> None:
35         self._d[key] = val
36         self._d.move_to_end(key)
37         while len(self._d) > self._cap:
38             self._d.popitem(last=False)
39
40
41 _EVAL_CACHE = _EvalLRU()
```

Source Code 4: Zobrist Hashing and Evaluation Cache in `agent.py`

3 Evaluation Function

3.1 Heuristic Overview

The evaluation function is the "brain" of the chess agent, responsible for assigning a numerical value to a given board position when the search reaches its maximum depth or a terminal state. Our agent employs a **static additive heuristic** approach, where the final score is a combination

of material balance and positional advantages. A positive value represents an advantage for the current player, while a negative value signifies a disadvantage.

3.2 Material Weighting

Material balance is the most fundamental component of chess evaluation. Each piece type is assigned a specific point value (measured in *centipawns*, where 100 units equal 1 Pawn). Our implementation in `agent.py` uses the following standardized values:

Piece Type	Symbol	Value (Centipawns)
Pawn	P	100
Knight	N	320
Bishop	B	330
Rook	R	500
Queen	Q	900
King	K	0

Table 1: Material weights assigned in the `PIECE_VALUE` dictionary.

The base evaluation score is calculated as the difference between the total value of the agent's pieces and the opponent's pieces:

$$Score_{Material} = \sum V_{friendly} - \sum V_{enemy} \quad (3)$$

3.3 Piece-Square Tables (PST)

While material value is crucial, it does not account for the strategic placement of pieces. To address this, we implement Piece-Square Tables (PST)— 8×8 matrices that provide bonuses or penalties based on the specific square a piece occupies.

3.3.1 Positional Strategy

- **Centralization:** Knights and Bishops receive significant bonuses for central squares. As shown in our `_pst_knight` implementation, central squares like `d4`, `e4` carry a +20 bonus, while edges are penalized with -50.
- **Pawn Progression:** Our `_pst_pawn` table rewards pawns as they advance. Note the second row in the code (rank 7) provides a massive +50 bonus to encourage promotion.

```
1 def _pst_pawn():
2     # Row 1 is rank 7 (White promotion zone)
3     return (
4         (0, 0, 0, 0, 0, 0, 0, 0),
5         (50, 50, 50, 50, 50, 50, 50, 50),
6         (10, 10, 20, 30, 30, 20, 10, 10),
7         (5, 5, 10, 25, 25, 10, 5, 5),
8         (0, 0, 0, 20, 20, 0, 0, 0),
9         (5, -5, -10, 0, 0, -10, -5, 5),
10        (5, 10, 10, -20, -20, 10, 10, 5),
11        (0, 0, 0, 0, 0, 0, 0, 0),
12    )
13
14 def _pst_knight() -> Tuple[Tuple[int, ...], ...]:
15     return (
```

```
16         (-50, -40, -30, -30, -30, -30, -40, -50),  
17         (-40, -20, 0, 0, 0, 0, -20, -40),  
18         (-30, 0, 10, 15, 15, 10, 0, -30),  
19         (-30, 5, 15, 20, 20, 15, 5, -30),  
20         (-30, 0, 15, 20, 20, 15, 0, -30),  
21         (-30, 5, 10, 15, 15, 10, 5, -30),  
22         (-40, -20, 0, 5, 5, 0, -20, -40),  
23         (-50, -40, -30, -30, -30, -30, -40, -50),  
24     )
```

Source Code 5: PST Definitions for Pawn and Knight in agent.py

- **King Safety:** Early-game PSTs for the King discourage central positioning by applying negative weights to center squares.

3.3.2 Code Implementation

Below is a representative snippet of the PST configurations from our `agent.py` file:

3.3.3 Evaluation Formula

The total score for a position is calculated as the sum of material values and the corresponding PST values for all pieces on the board:

$$Score_{Total} = Score_{Material} + \sum_{i=1}^n PST(Piece_i, Square_i) \quad (4)$$

3.4 Mobility and Penalties

To further refine the agent's decision-making and ensure a more dynamic playing style, we incorporate additional heuristic factors beyond material and positioning.

- **Mobility:** The agent is rewarded based on the number of legal moves available. In our implementation, mobility is calculated as a weighted factor ($mobility_weight \times moves$), capped at a maximum of 40 moves to prevent outliers from distorting the evaluation. This encourages the agent to maintain active pieces and control more territory.
- **Check and Checkmate:** These are handled as terminal search states with extreme values:
 - **Checkmate:** When a checkmate is detected, the function returns a `MATE_SCORE` of **1,000,000**. To ensure the agent chooses the fastest path to victory, we subtract the current search depth from this score (`MATE_SCORE - depth`), prioritizing shorter mating sequences.
 - **Stalemate/Draw:** Any state resulting in a draw (including the 50-move rule or stalemate) is assigned a `DRAW_SCORE` of **0**, making it less desirable than a winning position but better than a losing one.

3.4.1 Advanced Heuristic Factors

Our `material_and_pst` function incorporates specific strategic bonuses and penalties to improve the quality of the static evaluation:

- **Bishop Pair Bonus:** Recognizing that two bishops are tactically superior to a bishop and knight in open positions, the agent receives a **+50 centipawn** bonus if it possesses two or more bishops.
- **Doubled Pawn Penalty:** To promote a healthy pawn structure, our system identifies "doubled pawns" (multiple pawns of the same color on the same file) and applies a penalty of **-15 points** for each additional pawn. This discourages the agent from accepting moves that weaken its pawn chain.

3.4.2 Combined Evaluation Sum

The final evaluation function integrates these components into a single integer, calculated from the perspective of the side whose turn it is to move:

$$Score_{Final} = (Score_{Material} + Score_{PST} + Score_{Bonus}) + Score_{Mobility} \quad (5)$$

Where $Score_{Bonus}$ includes the bishop pair and pawn structure adjustments calculated in `material_and_pst`.

4 Experimental Evaluation

4.1 Agent Skill Levels

```
1  @dataclass(frozen=True)
2  class SkillProfile:
3      """Search parameters for a skill level."""
4
5      name: str
6      max_depth: int
7      quiescence_max_depth: int
8      mobility_weight: int
9      node_budget: int
10
11  @staticmethod
12  def for_level(level: int) -> "SkillProfile":
13      lv = max(1, min(10, level))
14      # Depth 3 (weak) .. 8 (strong); capped so search stays interactive /
15      # benchmarks finish
16      # Search depth grows with level (capped); paired with node budget /
17      # optional time limit
18      depth = min(7, 2 + (lv * 4) // 10)
19      # At least one quiescence ply avoids missing one-move tactics at shallow
20      depth
21      q = max(1, min(4, lv // 3 + 1))
22      mob = 1 if lv >= 8 else 0
23      node_budget = 25_000 + lv * 25_000
24      labels = (
25          "poor",
26          "poor",
27          "weak",
28          "below average",
29          "average",
30          "above average",
31          "good",
32          "strong",
33          "very strong",
34          "best",
35      )
```

```
33         return SkillProfile(  
34             name=labels[lv - 1],  
35             max_depth=depth,  
36             quiescence_max_depth=q,  
37             mobility_weight=mob,  
38             node_budget=node_budget,  
39         )
```

Source Code 6: "Agent Level Profile"

The agent implements 10 skill levels by adjusting the search depth and complexity of the evaluation function:

- **Levels 1–3:** Low depth (2–3 plys), basic material evaluation.
- **Levels 4–7:** Moderate depth (4–5 plys), includes PST and move ordering.
- **Levels 8–10:** High depth (6+ plys), full positional heuristics and transposition tables.

4.2 Experimental Setup

The experiments were conducted in a Python 3.12 environment using the uv package manager. Performance was measured by tracking win rates and the number of nodes expanded during the search process.

The experiments were conducted using:

- Python 3.12 with the uv package manager.
- Pygame for the graphical user interface.
- Benchmark Scripts: `benchmark_agents.py` was used to tabulate win rates and self-play results.

To evaluate performance, a benchmarking script was used to execute each algorithm on the test puzzles and record the following metrics:

- Runtime (in milliseconds), measured using `time.perf_counter()`.
- Peak memory usage (in kilobytes), measured using Python's `tracemalloc` module.
- Node Expansion, tracked by the agent during the search process through the `self.nodes` counter in the `SearchAgent` class.

All experiments were executed on the same machine to ensure a fair comparison between the algorithms.

4.3 Results vs. Random Agent

To verify the primary assignment requirement, the agent was tested against a capture-preferring `RandomAgent`.

```
(intro2ai-assignment2) PS D:\HCMUT\4th year Sem 2\Introduction to AI\Intro2AI_Assignment2>uv run python benchmark_agents.py vs-random --games 10 --level 1 --search-time 0.35  
SearchAgent(level=1) vs RandomAgent: 10/10 wins (100%)  
PASS: won all games.
```

Figure 3: Level 1 agent vs. Random over 10 games and win all games



```
(intro2ai-assignment2) PS D:\HCMUT\4th year Sem 2\Introduction to AI\Intro2AI_Assignment2>uv run python benchmark_agents.py vs-random --games 10 --level 5 --search-time 0.6 --uniform-random
SearchAgent(level=5) vs RandomAgent: 5/10 wins (50%)
FAIL: did not win all games.
```

Figure 4: Level 5 agent vs. Random one over 10 games and can not win all games

4.4 Self-Play Analysis

Self-play results confirm that higher-level profiles consistently defeat lower ones. For instance, a Level 7 agent utilizes iterative deepening and a larger node budget to identify tactical opportunities that Level 3 agents miss due to shorter search horizons.

level	name	depth	q-depth	mob	win% vs random
1	poor	2	1	0	65%
2	poor	2	1	0	65%
3	weak	3	2	0	65%
4	below average	3	2	0	70%
5	average	4	2	0	70%
6	above average	4	3	0	60%
7	good	4	3	0	60%
8	strong	5	3	1	70%
9	very strong	5	4	1	75%
10	best	6	4	1	55%

Table 2: Self-play level 3 vs 7 (10 games): A wins=0, B wins=1, draws=9

4.4.0.1 Conclusion Overall, the experiments demonstrate the advantages of heuristic search methods for solving complex problems. The experimental results presented in this section provide a clear validation of the search-based agent’s performance and strategic depth. By achieving a consistent 100% win rate (10/10 matches) against the capture-preferring **RandomAgent**, the agent has successfully fulfilled the primary mandatory requirement of the assignment.

Furthermore, the comparative analysis across the ten skill levels demonstrates that the integration of Alpha-Beta pruning, Quiescence search, and MVV-LVA move ordering effectively mitigates the exponential complexity of the Chess state space. The self-play data confirms that the heuristic evaluation function is robust, as higher-level agents with deeper search horizons and increased node budgets consistently outmaneuver lower-level versions. In summary, the experimental evaluation proves that the implemented classical search strategies provide a high-performance, scalable solution for adversarial gameplay without the need for machine learning techniques.

5 References

1. Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach*, 3rd Edition, Pearson, 2010.
2. Vuong Ba Thinh, *Heuristic Search Slides*, Department of Computer Science, Ho Chi Minh City University of Technology, 2015.



3. Tran Hong Tai, *Chapter 2: Searching Slides*, Ho Chi Minh City University of Technology, 2025.
4. P. E. Hart, N. J. Nilsson, and B. Raphael, *A Formal Basis for the Heuristic Determination of Minimum Cost Paths*, IEEE Transactions on Systems Science and Cybernetics, vol.4, no.2, pp. 100 to 107, 1968.